

[Home](#) [Syllabus](#) [Topics](#) [Homework](#) [Projects](#) [Resources](#) [FAQ](#) [Staff](#)

## Project 1: Llama Stacks

Due: Tuesday, October 4, 8:59:59pm

### Corrections and Addenda

- `pop()` should throw a `LlamaUnderflow` exception if the stack is empty.
- Add to `Llama.h`

```
#include <string>
```

### Objectives

The objective of this programming assignment is to have you review C++ programming using following features: object-oriented design, dynamic memory allocation, pointer manipulation, exceptions and templates.

### Introduction

The premise of this assignment is that you are using a system where calls to `new` and `delete` for dynamic memory allocation and de-allocation are very slow. They are slow enough that it is worth your while to consider data structures that reduce the number of calls to `new` and `delete`.

The data structure we are considering is a stack. This can be implemented as an array or as a linked list with advantages and disadvantages to both. However, your application does not allow you to determine the maximum size of the stack. If you implemented the stack using an array, you may run out of storage if the stack exceeds the size of the array. You can dynamically allocate a bigger array and copy the old stack, but the time for copying is prohibitively slow. Using a standard linked list would involve a call to `new` or `delete` for each push or pop instruction. This is also too slow.

Instead of a simple linked list, we will use a linked list where each node contains an array of items. The array will be of fixed size, say for example, 100. This linked list will start out with a single node that holds an array of 100 items. So, a single node can hold up to 100 items of the stack. That means we can push items on the stack and pop items off the stack without calling `new` or `delete` until the stack grows larger than 100. If that happens, we will add a node to the linked list. The result is a stack that holds 200 items. Adding space for 100 more items involved just one call to `new`.

There's a catch. Suppose our stack has 101 items and we do a pop. Now we only have 100 items, what do we do with the linked list? We can remove and deallocate the node that used to hold items 101 through 200. If we do that, what happens if we push right after the pop? We would have to add a node to the linked list again. If there's a sequence of push and pop instructions that causes the stack size to flip between 100 and 101, then we are left with the situation that every push and pop instruction results in a call to `new` or `delete`. Then there would be no advantage to our new data structure.

The solution is that we should not immediately remove a node when our stack size drops from 101 to 100. Instead, we should wait until the stack size drops much further, until it drops to 50 items (which requires at least 50 pop instructions to achieve). When the stack size is between 100 and 51, we will keep the "extra" node around in case the stack size grows to 101 again. Then we can use the space in the extra node and bypass the need to call new.

To recap, if our stack size reaches 101 for the first time, we add a node to the linked list. The stack size can increase and decrease, but we do not deallocate a node until the stack size drops to 50. Also, we won't have to allocate a new node until the stack size increases past 200. In general, we follow this scheme whenever a node is filled --- i.e., when the stack size grows past 200, 300, 400, ... This strategy prevents us from having to call new or delete very often. It also allows our stack to grow and shrink without having to copy the entire stack to a new location.

Let's call this data structure a Llama Stack (short for Linked-List Array Mixed Amalgamated Stack).

## Assignment

Your assignment is to implement *and test* a templated Llama Stack in C++. You should use the following two header files:

```
#ifndef _LLAMANODE_H_
#define _LLAMANODE_H_

/* File: LlamaNode.h

UMBC CMSC 341 Fall 2016 Project 1

This file has the class declaration for the LlamaNode class
for Project 1. See project description for details.

This file should not be modified in ANY way.

When your program is graded, it will be compiled with the original
version of this file. Your program must work with the original.

*/

#include <iostream>
using namespace std ;

template <class T, int LN_SIZE>    // forward class declaration
class Llama ;

template <class T, int LN_SIZE>
class LlamaNode {

    friend class Llama<T,LN_SIZE> ;

public:

    LlamaNode() ;
```

```

~LlamaNode() ;

static void report() ;

private:

static int newCount ;    // # of times constructor was called
static int deleteCount ; // # of times desctructor was called

T arr[LN_SIZE] ;
LlamaNode *m_next ;

} ;

#include "LlamaNode.cpp"

#endif

```

```

#ifndef _LLAMA_H_
#define _LLAMA_H_

/* File: Llama.h

UMBC CMSC 341 Fall 2016 Project 1

This file has the class declaration for the LlamaNode class
for Project 1. See project description for details.

You may add public and private data members to the Llama class.

You may add public and private member functions to the Llama class.

*/

#include <stdexcept>
#include "LlamaNode.h"

using namespace std ;

class LlamaUnderflow : public std::out_of_range {

public:

    LlamaUnderflow(const string& what) : std::out_of_range(what) { }

} ;

```

```

template <class T, int LN_SIZE>
class Llama {

    public:

        Llama() ;
        Llama(const Llama<T,LN_SIZE>& other) ;    // copy constructor
        ~Llama() ;

        int size() ;
        void dump() ;
        void push(const T& data) ;
        T pop() ;

        void dup() ;    // (top) A B C D -> A A B C D
        void swap() ;    // (top) A B C D -> B A C D
        void rot() ;    // (top) A B C D -> C A B D

        T peek(int offset) const ;

        // overloaded assignment operator
        //
        const Llama<T,LN_SIZE>& operator=(const Llama<T,LN_SIZE>& rhs) ;

        //
        // Add your public member functions & public data mebers here:
        //

    private:

        //
        // Add your private member functions & private data mebers here:
        //

};

#include "Llama.cpp"

#endif

```

The implementation of the `LlamaNode` class is trivial and has been done for you. (See [LlamaNode.cpp](#).) In fact, you are not allowed to change the `LlamaNode` files in any way.

Your assignment is to implement the member functions of the `Llama` class. You will need to add data members and member functions to the `Llama` class, but the rest of the header file should not be changed. You must implement the following member functions:

- `Llama() ;`

This is the default constructor for the `Llama` class. Since you are picking the data members for this class, much of this code is determined by you. At the very least, the default constructor must create a `Llama` Stack that holds one node.

- `Llama(const Llama<T, LN_SIZE>& other) ;`

This is the copy constructor for the `Llama` class. It should make a complete copy of the `Llama` Stack given in the parameter. The target of the copy is the host object.

You should not call the assignment operator from the copy constructor (or vice versa). The objective of these two member functions are sufficiently different that you should implement them separately. If you do not want to duplicate similar code in the two functions, make a third function that you can call from the copy constructor and from the assignment operator.

- `~Llama() ;`

This is the destructor for the `Llama` class. All dynamically allocated memory associated with the host object must be deallocated. You should use `valgrind` on GL to check for memory leaks.

- `int size() ;`

This function returns the number of items in the `Llama` Stack.

- `void dump() ;`

This member function is used for debugging. It should print out all pertinent information regarding the host `Llama` stack to `cerr`, including the number of `LlamaNodes` that have ever been created and destroyed. It should also print out the address of each node and whether there is currently an "extra" node in the data structure.

See the sample outputs below for the suggested format of the output of the `dump()` function.

- `void push(const T& data) ;`

The `push()` member function adds `data` to the top of the stack. A copy of `data` should be made.

- ```
T pop() ;
```

The `pop()` member function removes and returns the item at the top of the Llama Stack.

If the stack is empty, `pop()` should throw a `LlamaUnderflow` exception. (The `LlamaUnderflow` class is defined in [Llama.h](#).)

- ```
void dup() ;
```

The `dup()` member function duplicates the top of the stack. For example, suppose the stack originally held A B C D (where A is at the top). Then, after calling `dup()`, the stack should become A A B C D.

If the stack is empty, `dup()` should throw a `LlamaUnderflow` exception.

- ```
void swap() ;
```

The `swap()` member function exchanges the top two items of the stack. For example, suppose the stack originally held A B C D (where A is at the top). Then, after calling `swap()`, the stack should become B A C D.

If there are fewer than two items in the stack, `swap()` should throw a `LlamaUnderflow` exception.

- ```
void rot() ;
```

The `rot()` member function permutes the top three items of the stack. For example, suppose the stack originally held A B C D (where A is at the top). Then, after calling `rot()`, the stack should become C A B D.

Note: "rot" is short for "rotate".

If there are fewer than three items in the stack, `rot()` should throw a `LlamaUnderflow` exception.

- ```
T peek(int offset) const ;
```

The `peek()` member function returns the value of an item in the stack. The offset is used to determine which item is retrieved: `peek(0)` is the top of the stack, `peek(1)` is the first item below the top, ... The stack should not be modified in any way after a call to `peek()`. If the offset is too large, `peek()` should throw a `LlamaUnderflow` exception.

- ```
const Llama<T, LN_SIZE>& operator=(const Llama<T, LN_SIZE>& rhs) ;
```

This is the overloaded assignment operator. The assignment operator should handle the case of self-assignment. Any existing dynamically allocated memory in the host object must be deallocated beforehand. A complete copy of the right

hand side of the assignment should be placed in the host object. This should handle the possible existence of an "extra" node in rhs.

You should not call the assignment operator from the copy constructor (or vice versa). See note in copy constructor description.

Finally, you must write a `main()` function that should be placed in a file named `Driver.cpp`. Your `main()` function should thoroughly test your implementation.

## Files and Sample Runs

Here are some files for you to download:

- [LlamaNode.h](#) (do not change!)
- [LlamaNode.cpp](#) (do not change!)
- [Llama.h](#)

The following three test programs may be used to check the compatibility of your implementation. These programs do not check the *correctness* of your implementation. Even if your implementation compiles and runs correctly with these programs, it does not mean your implementation is error-free. Grading will be done using programs that exercise your implementation much more thoroughly. You must do the testing yourself --- testing is part of programming. Conversely, if your implementation does *not* compile or does *not* run correctly with these test programs, then it is unlikely that it will compile or run correctly with the grading programs.

- [driver1.cpp](#): uses Llama Stack with strings.  
Output from sample run: [driver1.txt](#).  
Output from sample run with valgrind: [driver1v.txt](#).
- [driver2.cpp](#): uses Llama Stack with `float`. Also catches an exception.  
Output from sample run: [driver2.txt](#).  
Output from sample run with valgrind: [driver2v.txt](#).
- [driver3.cpp](#): uses Llama Stack with a class that has dynamically allocated data.  
Compile with [OvCoInt.h](#) and [OvCoInt.cpp](#) (Note: OvCoInt = "Overly Complicated Integer").  
Output from sample run: [driver3.txt](#).  
Output from sample run with valgrind: [driver3v.txt](#).

## Implementation Notes

- Yes, you can have `int` template parameters.
- Do not put a `main()` function in `Llama.cpp`. Your `main()` function should go in a separate file called `Driver.cpp`.
- Remember that your `Llama.cpp` file must also be guarded. This is because the `Llama.h` file includes `Llama.cpp` (and `Llama.cpp` also includes `Llama.h`). This arrangement is necessary for C++ templates.
- Since a `LlamaNode` can be quite large, your linked list should not use a dummy header node. They are not very useful for stacks anyway.
- Except for `dump()`, `peek()`, the copy constructor, the destructor and the assignment operator, all member functions should run in  $O(1)$  time. I.e., the running time for the member function should not depend on the size of the stack. (We will consider `LN_SIZE` to be a constant.)
- You will need a consistent design philosophy. You will need data members that determine where the top of the stack is, how many items there are in the stack, whether there is currently an "extra" node... You should think a lot about how,

when and where these data members are updated and kept consistent.

- There should be at most one extra node in the data structure at any time. (I.e., you can't just never delete `LlamaNodes`.)
- Recommendation: Use the C/C++ `assert` macro facility to test your design. To do this:

```
#include <assert.h>
```

In your program, you can use `assert()` with a Boolean expression that "should" be true. For example, if you think that a pointer, `ptr`, should always be non-NULL at a particular point in your code, then you can use:

```
assert( ptr != NULL ) ;
```

If `ptr` is indeed not NULL, nothing happens. On the other hand, if `ptr` is NULL, then the program will stop and print out the line number where the assertion failed.

For example, the third sample driver above, [driver3.cpp](#), uses the Overly Complicated Integer class (see [OvCoInt.h](#) and [OvCoInt.cpp](#)) which has a dynamically allocated int. The pointer to this location is never supposed to be NULL. The constructor for `OvCoInt` allocates space and it *shouldn't* ever be deallocated until the destructor is invoked. The `assert()` macro is used throughout to check.

The `assert()` macro is a debugging tool used during development. It is also an example of defensive programming. After programs are fully debugged, the `assert()` macro can be turned off by compiling with the `-DNDEBUG` flag. So, you do not need to worry about efficiency and there is no need to comment out all the `assert()` calls when you are done debugging.

- Do not implement `dump()` using `peek()`. The `dump()` function should be simple and bug free. Use `dump()` to debug more complicated functions like `peek()`.
- Approach this assignment using incremental development. Implement `push()` and `dump()` first. Make sure that `push()` and `dump()` are working. Then implement `pop()`. Make sure you are handling the "extra" node correctly before working on the copy constructor and the assignment operator.
- Along the lines of incremental development, try working without templates first, because non-templated code is easier to debug. Remove or comment out the template syntax and add the following macros to

`Llama.h`

```
#define T string  
#define LN_SIZE 4
```

After that, the identifiers `T` and `LN_SIZE` will be replaced with "string" and "4" respectively. After verifying your code works with strings, check if it works with `int`. Then, check it with `OvCoInt`. (Just change the `#defines`.) Finally, check if it works with different values of `LN_SIZE`. After that, making templates from your code is mostly a typing exercise.

- Test your programs for memory leaks (dynamically allocated memory that was never released) using `valgrind`. The `valgrind` command is available on GL. Just compile your test program and run:

```
valgrind ./Driver.out
```

This is assuming that your executable file is named `Driver.out`. If that run did not leak memory, the output from `valgrind` will say:

```
All heap blocks were freed -- no leaks are possible
```

As usual, the fact that a single run of your implementation did not leak memory does not mean that it will never leak memory.



## What to Submit

You must submit the following files to the `proj1/src` directory.

- `Llama.h`
- `Llama.cpp`
- `Driver.cpp`

You do not need to submit `LlamaNode.h` and `LlamaNode.cpp` because those files should not have changed. If you do happen to place a copy of `LlamaNode.h` and/or `LlamaNode.cpp` in your submission directory, they will be replaced by a copy of the original version.

If you followed the instructions in the [Project Submission](#) page to set up your directories, you can submit your code using these Unix commands.

```
mkdir ~/cs341proj/proj1/src
cp Llama.h Llama.cpp Driver.cpp ~/cs341proj/proj1/src/
```

Note: you only have to use `mkdir` to make the `src` subdirectory once. If you submit again, just use the `cp` command.

## Discussion Topics

Here are some topics to think about:

- What if `LN_SIZE` is 1?
- What if `LN_SIZE` is bigger than the number of items you will ever store in the stack?
- The `LlamaNode` class does not have a copy constructor or an overloaded assignment operator. Why is this OK?